

# Deal with big data in R using bigmemory package

Xiaojuan Hao

Department of Statistics  
University of Nebraska-Lincoln  
April 28, 2015

# ➤ Background

## ❖ What Is Big Data

- Size (We focus on)
- Complexity
- Rate of growth

## ❖ Problem In R For The Big Size Of Data

- Size of 2 GB
- Runs on a windows 7 system with 8 GB of RAM.
- Takes about 9 minutes with `read.table` function
- Consume 6 GB of memory.
- Some simple calculations will cause error:

Error message : Reached total allocation of 7987Mb: see  
`help(memory.size)`

```

> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 183080  9.8      407500 21.8      183080  9.8
Vcells 278241  2.2      786432  6.0      278241  2.2
> start.time<-proc.time()
> xx <- read.table("bigdata.txt", header = T)
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:", save.time[3]/60, "\n \n")

```

```

Number of minutes running: 9.269833

```

```

> gc()
      used      (Mb) gc trigger      (Mb) max used      (Mb)
Ncells   186123   10.0      686864   36.7      666489   35.6
Vcells 248160566 1893.4  702365274 5358.7 857537920 6542.5
>
> xx+xx
Error: cannot allocate vector of size 377.9 Mb
In addition: Warning messages:
1: In structure(list(message = as.character(message), call = call),  :
  Reached total allocation of 7987Mb: see help(memory.size)

```

1. You can't load the data in R with only 4 GB of memory.
2. The gc function is used to monitor the usage of memory .
3. The element in the last column and second rows shows the maximum space used in Mb since the last call to gc (reset = TRUE)

## ❖ About R

### ❑ We love R because

- Flexible for data analysis
- Big graphical capabilities
- Extensible
- Free and available for different platforms.

### ❑ Two major limitations

- Uses only 1 core by default on CPU
- R reads data into memory by default
  1. Exhaust RAM by storing unnecessary data.
  2. More memory will lead system to be frozen.

## ❖ Couple of Solutions

- Buy more RAM

Expensive

- Apply more memory on HCC

Not convenient

- Use C/C++ program

Lack the flexibility and convenience of R's rich environment.

- Use R packages for big data support

1. bigmemory      2. ff

- Use parallelism to process and generate large data sets on cluster

1. Hadoop              2. MapReduce

# ➤ **Outline**

## ❖ **Introduce bigmemory Package**

- Why we need bigmemory
- About the big family

## ❖ **How To Use bigmemory Package**

- Some main functions.
- Shared memory with foreach iteration
- Use biglm package with bigmemory
- Other commands

## ❖ **Summary of The bigmemory Package**

# ➤ **Introduce bigmemory**

## ❖ **What Is bigmemory**

- Manage massive matrices with shared memory and memory-mapped file

## ❖ **Why We Need bigmemory**

- Store big matrices in memory and support their basic manipulation and exploration. This procession was managed in R but implemented in C++
- Very simple to use
- Multiple processors on the same machine can access to the same big data sets by shared memory.

## ❖ Big Family

The big family contains several packages for analysis of big datasets

- **bigmemory**

Provides the core matrix-like support.

Functions include: `nrow`, `ncol`, `dim`, `tail`, `head`, `apply`, `big.matrix`,  
`read.big.matrix`, `mwhich`.

- **biganalytics**

Provides routine analysis on big matrix.

Functions include: `sum`, `range`, `mean`, `colsum`, `colrange`, `colmean`,  
`biglm.big.matrix`, `bigglm.big.matrix`

- **bigtabulate**

Adds table and split-like support for `big.matrix` objects

- **bigalgebra**

Provides linear algebra operations on `R matrix` as well as `big.matrix`.



# ➤ Some Main Functions

## ❖ Read In Data Functions

```
read.big.matrix(filename, sep = ',', header = FALSE,  
                col.names = NULL, row.names = NULL,  
                type = NA, skip = 0, separated = FALSE,  
                backingfile = NULL, backingpath = NULL,  
                descriptorfile = NULL, extraCols = NULL,  
                shared=TRUE)
```

type: integer (4 bytes), short(2 bytes), double(8 bytes), char(1 bytes)

backingfile: the root name for the file(s) for the cache of x.

backingpath: the path to the directory containing the backingfile

descriptorfile: the name of the file to hold the backingfile description.

shared: if TRUE, the `big.matrix` can be shared across processes.

# Example:

## Step 1: Simulate one data set

```
mydata=matrix(c(NA),nrow=10072112,ncol=5)
set.seed(12345)
mydata[,1]=sample(c(1:17770), 10072112, replace = TRUE)
mydata[,2]=sample(c(1:480189), 10072112, replace = TRUE)
mydata[,3]=sample(c(1:5), 10072112, replace = TRUE)
mydata[,4]=sample(c(1999:2005), 10072112, replace = TRUE)
mydata[,5]=sample(c(1:12), 10072112, replace = TRUE)
write.table(mydata, file = "example.txt", sep = " ",row.names
= F, col.names = F)
```

- 10 million rows and 5 columns
- File has size of 215.1055 MB

## Step 2: Read in the “example.txt” file

### 1. Use usual R read.table function

```
> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 183080  9.8      407500 21.8    183080  9.8
Vcells 278241  2.2      786432  6.0    278241  2.2
> start.time<-proc.time()
> x <- read.table("example.txt",colClasses = "integer", header=F,
      col.names = c("movie", "customer", "rating","year",
                    "month"))
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:", save.time[3]/60, "\n \n")
Number of minutes running: 0.6815
> gc()
      used (Mb) gc trigger (Mb) max used (Mb)
N 186106  10.0      407500  21.8    188080  10.1
V 25463925 194.3    55646223 424.6  51920931 396.2
> dim(x)
[1] 10072112      5
```

## 2. Use read.big.matrix function

```
> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 287878 15.4      467875   25    287878 15.4
Vcells 429260  3.3      905753    7    429260  3.3
> start.time<-proc.time()

➤ x <- read.big.matrix("example.txt", header =F,type = "integer",sep = "
      ",backingfile ="data.bin", descriptor = "data.desc",col.names =
      c("movie", "customer","rating","year", "month"), shared=TRUE)
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:",save.time[3]/60, "\n \n")
Number of minutes running: 0.8

> gc()
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 290574 15.6      531268 28.4    303567 16.3
Vcells 432636  3.4      905753  7.0    905084  7.0
> dim(x)
[1] 10072112      5
```

a) Spend more time

b) Save memory

- c) `read.big.matrix` function creates the binary file-backing associated with the `big.matrix` object `x` and shared descriptor file.
- d) Use shared descriptor, we can load the data from disk with `attach.big.matrix` function.

```
> start.time<-proc.time()
> datadesc<-dget("data.desc")
> data<-attach.big.matrix(datadesc)
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:", save.time[3]/60, "\n \n")
  Number of minutes running: 0.0006666667
> data
An object of class "big.matrix"
Slot "address":
<pointer: 0x000000000001c9ec0>
> head(data)
      movie customer rating year month
[1,] 12811    121761      5 2004     12
[2,] 15563    351198      2 2005      9
[3,] 13523    173425      5 2002      3
[4,] 15747    132685      1 2003      4
[5,]  8112    401461      4 2003      9
[6,]  2957    346798      3 2000      3
```

e) `big.matrix` holds an pointer to a C++ matrix that is on disk

```
> x
An object of class "big.matrix"
Slot "address":
<pointer: 0x000000000003496b0>

> head(x)
      movie customer rating year month
[1,] 12811    121761      5 2004     12
[2,] 15563    351198      2 2005      9
[3,] 13523    173425      5 2002      3
[4,] 15747    132685      1 2003      4
[5,]  8112    401461      4 2003      9
[6,]  2957    346798      3 2000      3

> is.filebacked(x)
[1] TRUE
> is.big.matrix(x)
[1] TRUE
> is.shared(x)
[1] TRUE
```

```
> datadesc<-dget("data.desc")
> data<-attach.big.matrix(datadesc)
> data
An object of class "big.matrix"
Slot "address":
<pointer: 0x000000000001c9ec0>

> head(data)
      movie customer rating year month
[1,] 12811    121761      5 2004     12
[2,] 15563    351198      2 2005      9
[3,] 13523    173425      5 2002      3
[4,] 15747    132685      1 2003      4
[5,]  8112    401461      4 2003      9
[6,]  2957    346798      3 2000      3

> is.filebacked(x)
[1] TRUE
> is.big.matrix(x)
[1] TRUE
> is.shared(x)
[1] TRUE
```

- f) The matrices can contain only numeric (char, short, int, double) values.
- g) Need preprocess data so that factors are coded as numeric values.

Example: Read in test.csv file

- i. Total 20 patients
- ii. Three columns. (PID (integer), BP(factor), W (integer))

```
> test <- read.big.matrix("test.csv", header = T, type = "integer")
> test
An object of class "big.matrix"
Slot "address":
<pointer: 0x0000000019bcd100>
```

```
> head(test[,])
      PID BP    W
[1,]   1 NA   90
[2,]   2 NA  138
[3,]   3 NA  170
[4,]   4 NA  112
[5,]   5 NA  130
[6,]   6 NA  100
```

## ❖ Construct Matrix Functions

```
big.matrix(nrow, ncol, type, init = NULL, dimnames = NULL,  
           separated = FALSE, backingfile = NULL,  
           backingpath = NULL, descriptorfile = NULL,  
           binarydescriptor=FALSE, shared = TRUE)
```

**Example 1**

```
>a<- big.matrix(10, 2, type='integer', init=-5)  
> options(bigmemory.allow.dimnames=TRUE)  
> colnames(a) <- c("alpha", "beta")  
> is.big.matrix(a)  
[1] TRUE  
> is.shared(a)  
[1] TRUE  
> is.filebacked(a)  
[1] FALSE
```

**Example 2**

```
> z <- filebacked.big.matrix(3, 3, type='integer', init=123,  
                             backingfile="example.bin", descriptorfile="example.desc",  
                             dimnames=list(c('a','b','c'), c('d', 'e', 'f')))  
> is.big.matrix(z)  
[1] TRUE  
> is.shared(z)  
[1] TRUE  
> is.filebacked(z)  
[1] TRUE
```



## ■ Several Matrix Objects

### 1. `big.matrix`

- a) Points to a data structure in C++.
- b) For a single R process.
- c) Limited by available RAM

### 2. `shared.big.matrix`

- a) Similar to `Big.matrix`
- b) Can be shared among multiple R processes

### 3. `filebacked.big.matrix`

- a) Points to a file on disk containing the matrix
- b) Files can be shared across a cluster

## ❖ mwhich Function

- Provides efficient row selections for `big.matrix` and `matrix` objects.
- Based loosely on R's `which` function without memory overhead

```
> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 292363 15.7      531268 28.4    292363 15.7
Vcells 436893  3.4      905753  7.0    436893  3.4
> cust.indices.inefficient <- which(data[, "customer"] == as.integer(6))
> gc()
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 293247 15.7      531268 28.4    303777 16.3
Vcells 437744  3.4     9539607 72.8 10669871 81.5
> head(data[cust.indices.inefficient,])
      movie customer rating year month
[1,]   508         6      1 2004     4
[2,]  2652         6      2 1999     8
[3,]  1171         6      2 2001    10
[4,] 15662         6      2 2001     7
[5,]   7139         6      5 2001     7
[6,]   858         6      3 2000     5
```

**`mwhich(x, cols, vals, comps, op = 'AND')`**

**x:** a `big.matrix` object.

**cols:** a vector of column indices or names.

**vals:** a list of vectors of length 1 or 2; length 1. It is used to test equality (or inequality),

**comps:** a list of operators including 'eq', 'neq', 'le', 'lt', 'ge' and 'gt'.

**op:** either 'AND' or 'OR'.

```
> gc(reset=T)
```

|        | used (Mb)   | gc trigger (Mb) | max used (Mb) |
|--------|-------------|-----------------|---------------|
| Ncells | 293350 15.7 | 531268 28.4     | 293350 15.7   |
| Vcells | 437916 3.4  | 7631685 58.3    | 437916 3.4    |

```
> cust.indices <- mwhich(data, "customer", 6, "eq")
```

```
> gc()
```

|        | used (Mb)   | gc trigger (Mb) | max used (Mb) |
|--------|-------------|-----------------|---------------|
| Ncells | 294316 15.8 | 531268 28.4     | 301658 16.2   |
| Vcells | 438252 3.4  | 6105348 46.6    | 771162 5.9    |

```
> sum(cust.indices.inefficient != cust.indices)
```

```
[1] 0
```

```
> head(data[cust.indices, ])
```

|      | movie | customer | rating | year | month |
|------|-------|----------|--------|------|-------|
| [1,] | 508   | 6        | 1      | 2004 | 4     |
| [2,] | 2652  | 6        | 2      | 1999 | 8     |
| [3,] | 1171  | 6        | 2      | 2001 | 10    |
| [4,] | 15662 | 6        | 2      | 2001 | 7     |
| [5,] | 7139  | 6        | 5      | 2001 | 7     |
| [6,] | 858   | 6        | 3      | 2000 | 5     |

```
> these <- mwhich(data, c("customer", "rating"), list(6, 2),  
                  list("eq", "le"), "AND")
```

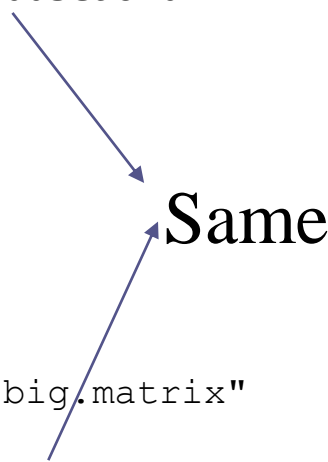
```
> head(data[these, ])
```

|      | movie | customer | rating | year | month |
|------|-------|----------|--------|------|-------|
| [1,] | 508   | 6        | 1      | 2004 | 4     |
| [2,] | 2652  | 6        | 2      | 1999 | 8     |
| [3,] | 1171  | 6        | 2      | 2001 | 10    |
| [4,] | 15662 | 6        | 2      | 2001 | 7     |
| [5,] | 4035  | 6        | 2      | 2002 | 4     |
| [6,] | 1970  | 6        | 1      | 1999 | 10    |

## ❖ deepcopy Function

```
deepcopy(x, cols = NULL, rows = NULL, y = NULL,  
         type = NULL, separated = NULL, backingfile =  
         NULL, backingpath = NULL, descriptorfile =  
         NULL, binarydescriptor=FALSE, shared=TRUE)
```

```
> z  
An object of class "big.matrix"  
Slot "address":  
<pointer: 0x000000000003c9820>  
  
> z[,]  
      d     e     f  
a 123 123 123  
b 123 123 123  
c 123 123 123  
  
> y<-z  
> y  
An object of class "big.matrix"  
Slot "address":  
<pointer: 0x000000000003c9820>
```



Same

```
> y[,]  
      d     e     f  
a 123 123 123  
b 123 123 123  
c 123 123 123  
> y[1,1]<-as.integer(1)  
> z[,]  
      d     e     f  
a   1 123 123  
b 123 123 123  
c 123 123 123  
  
> w<-deepcopy(z)  
> w  
An object of class "big.matrix"  
Slot "address":  
<pointer: 0x000000000002cd2a0>
```

z and y point to the same data in memory.

## ❖ flush Function

- `flush()` forces any modified information to be written to the file-backing.

```
> library(bigmemory)
> w <- read.big.matrix("w.txt", header
  = F, type = "integer", sep = " ",
  backingfile = "w.bin",
  descriptor = "w.desc",
  col.names = c("a", "b", "c"))
> w[, ]
      a    b    c
[1,] 123 123 123
[2,] 123 123 123
> w[,1] <- -1
> w[, ]
      a    b    c
[1,]  1 123 123
[2,]  1 123 123
> flush(w)
> neww = attach.big.matrix(dget("w.desc"))
> neww[, ]
      a    b    c
[1,]  1 123 123
[2,]  1 123 123
```



The backed file is changed

```
> w=read.big.matrix("w.txt", header = F, type = "integer",sep = " ")
> w[,]
      [,1] [,2] [,3]
[1,]  123  123  123
[2,]  123  123  123

> neww=attach.big.matrix(dget("w.desc"))
> neww[,]
      a    b    c
[1,]  1 123 123
[2,]  1 123 123
```

- It only changes the backed file.
- The original data set does not be changed
- The backed file does not match the original data anymore !!!

# ➤ Shared Memory

## ❖ Interactive Shared Memory

### Example 1

1. Two R sessions are connected to the same `shared.big.matrix`
2. The assignment in one process will affect the value in the other sessions.


#### Session 1

```
> r<-mwhich(x,"movie",2845,"eq")
> mean(x[r,"rating"])
[1] 3
> sd(x[r,"rating"])
[1] 1.406684

> mean(x[r,"rating"])
[1] 100
> sd(x[r,"rating"])
[1] 0
```

#### Session 2

```
> options(bigmemory.typecast.warning=FALSE)
> r<-mwhich(data,"movie",2845,"eq")
> data[r,"rating"]<-100
> mean(data[r,"rating"])
[1] 100
> sd(data[r,"rating"])
[1] 0
```





## ❖ Parallel Computation With bigmemory

### Example 2

1. Calculate the average rate for each movie.
2. Takes about 2.5 hours using for loop iteration.

```
> start.time<-proc.time()
> movie_uniq<-data[unique(data[,1]),1]
> n<-length(movie_uniq)
> movie_av_rate<-big.matrix(n,2, type='double')
> movie_av_rate[,1]<-movie_uniq
> for (i in 1:n)
  movie_av_rate[i,2] <-
    mean(data[mwhich(data,"movie",movie_uniq[i],"eq"),"rating"])
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:", save.time[3]/60, "\n \n")
```

Number of minutes running: 137.1643

### 3. Takes about 1.5 hours for parallel computation with two cores

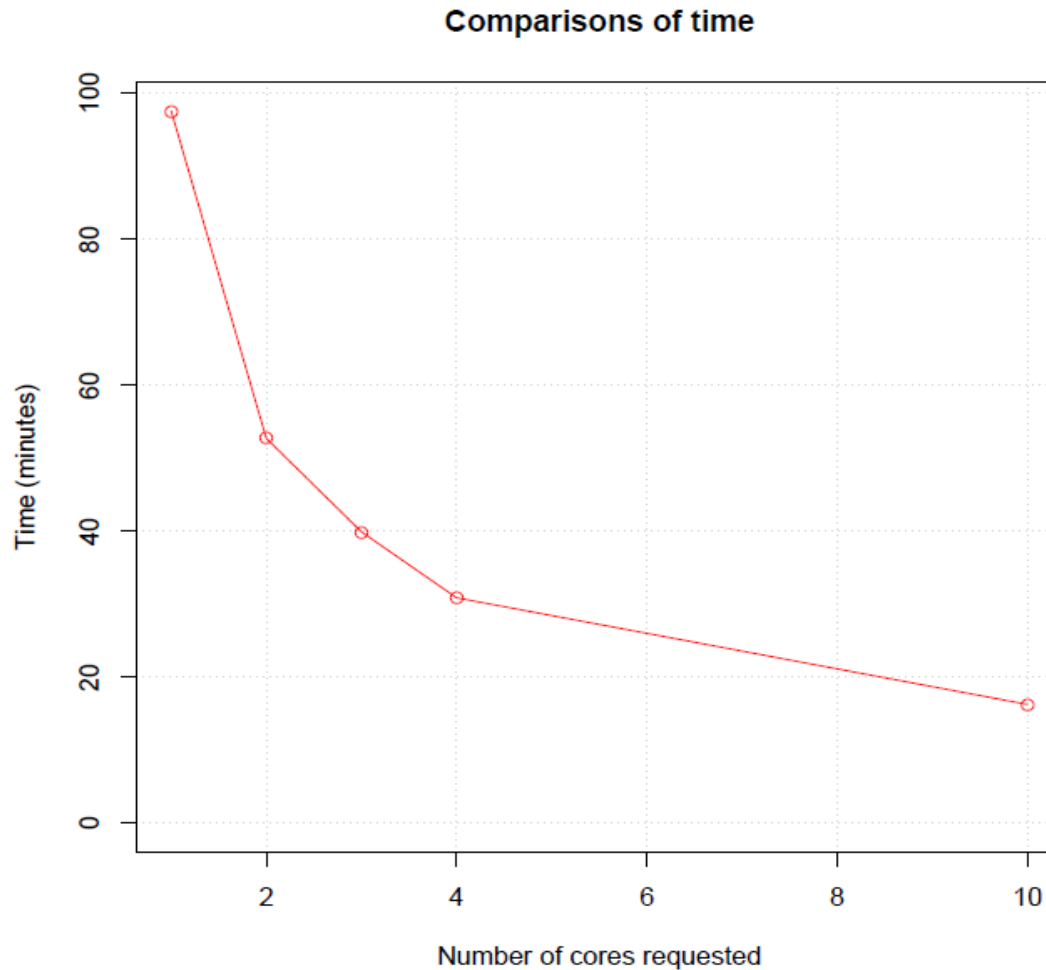
```
> library(bigmemory)
> datadesc<-dget("data.desc")
> data<-attach.big.matrix(datadesc)
> movie_uniq<-data[unique(data[,1]),1]
> n<-length(movie_uniq)
> movie_av_rate<-big.matrix(n,2, type='double')
> movie_av_rate[,1]<-movie_uniq
>
> library(doParallel)
> cl<-makeCluster(spec = 2)
> registerDoParallel(cl = cl)
> library(foreach)
> start.time<-proc.time()
> clusterSetRNGStream(cl = cl, iseed = 9182)
> res<-foreach(i = 1:n,.combine = rbind) %dopar% {
  require(bigmemory)
  data<-attach.big.matrix(datadesc)
  mean(data[mwhich(data,"movie",movie_uniq[i],"eq"),"rating"])
}
> stopCluster(cl)
> movie_av_rate[,2]<-res
> end.time<-proc.time()
> save.time<-end.time-start.time
> cat("\n Number of minutes running:", save.time[3]/60, "\n\n")
  Number of minutes running: 81.845
```

**Takes about 1.5 hours**



## 4. Use Crane

Time in minutes vs. the number of threads



## ➤ **biglm Package With bigmemory**

- `biganalytics` provides linear and generalized linear models for big data based on `biglm` package.
- `biglm.big.matrix` and `bigglm.big.matrix` save memory than usual `lm` function for big data regression.

### **Example**

For movie data, we try to predict the customer ratings using the factor of movie released year

```
library(biglm)
# library(biganalytics)
lm.b<-biglm.big.matrix(rating ~ year, data=x, fc="year")
```

```
> summary(lm.b)
```

```
Large data regression model: biglm(formula = formula, data = data, ...)
```

```
Sample size = 10072112
```

|             | Coef    | (95% CI)       | SE     | p      |
|-------------|---------|----------------|--------|--------|
| (Intercept) | 3.0072  | 3.0045 3.0098  | 0.0013 | 0.0000 |
| year2000    | -0.0024 | -0.0061 0.0014 | 0.0019 | 0.2069 |
| year2001    | -0.0016 | -0.0054 0.0021 | 0.0019 | 0.3824 |
| year2002    | -0.0027 | -0.0065 0.0010 | 0.0019 | 0.1496 |
| year2003    | 0.0001  | -0.0036 0.0039 | 0.0019 | 0.9489 |
| year2004    | -0.0006 | -0.0043 0.0032 | 0.0019 | 0.7643 |
| year2005    | -0.0009 | -0.0046 0.0029 | 0.0019 | 0.6495 |

$$\widehat{Rating} = \hat{a} + \hat{b}_1 * year2000 + \hat{b}_2 * year2001 + \hat{b}_3 * year2002 + \hat{b}_4 * year2003 + \hat{b}_5 * year2004 + \hat{b}_6 * year2005$$

- It appears that there is no linear relationship between the movie ratings and released years

## ➤ Other Commands

- Write a `big.matrix` to a file using `write.big.matrix` function.
- Parallel `foreach` function with `bigmemory` package will be the most efficient way to deal with big data sets.
- Save memory using correct type.

```
> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 295907 15.9      531268 28.4   295907 15.9
Vcells 440279  3.4     4884278 37.3   440279  3.4
> data[,4]<-data[,4]+1
> gc()
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells  296361 15.9      531268 28.4   301473 16.2
Vcells 10512493 80.3   30178098 230.3 35852849 273.6
> gc(reset=T)
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells 296369 15.9      531268 28.4   296369 15.9
Vcells 440408  3.4    24142478 184.2   440408  3.4
> data[,4]<-data[,4]-as.integer(1)
> gc()
      used (Mb) gc trigger (Mb) max used (Mb)
Ncells  296361 15.9      531268 28.4   299041 16.0
Vcells 5476437 41.8    19313982 147.4 20745705 158.3
```

## ➤ Summary of bigmemory

- Store a matrix in memory and easy to access without reload the data.
- Share matrices among multiple R sessions and clusters.
- Be careful when you use shared matrix.
- More efficient to use parallel calculation with bigmemory package for large data sets.
- Only deal with numeric matrices.
- More big family and functions are available.

See <http://cran.r-project.org/web/packages/bigmemory/bigmemory.pdf>

<http://cran.r-project.org/web/packages/bigmemory/index.html>

## ➤ Reference

- The Bigmemory Project, <http://www.bigmemory.org/>, the home of R packages bigmemory, biganalytics, bigtabulate, bigalgebra, and synchronicity. Packages available from CRAN or R-Forge.
- Emerson JW, Kane MJ (2009). “The R Package bigmemory: Supporting Efficient Computation and Concurrent Programming with Large Data Sets.” Journal of Statistical Software, Volume VV, Issue II.
- 2009 JSM Data Expo: Airline on-time performance. <http://stat-computing.org/dataexpo/2009/>.
- “The Bigmemory Project” by Michael Kane and John Emerson: April 29, 2010. <http://cran.r-project.org/web/packages/bigmemory/vignettes/Overview.pdf>
- “Taking R to the limit” by Ryan R Rosario: August 17, 2010. [http://www.bytemining.com/wp-content/uploads/2010/08/r\\_hpc\\_II.pdf](http://www.bytemining.com/wp-content/uploads/2010/08/r_hpc_II.pdf)